

Modul 215: Machine Learning

Linear Regression Methods: Linear & Logistic Regression

Wintersemester 2023/24

Tobias Kortus

Center for Technology and Transfer (ZTT)
University of Applied Sciences

November 9, 2023




Recap: Regression

- **Goal:** Find a function f (in most cases a set of parameters) to predict a quantity y for input x by using training data:

$$\left((x^{(1)}, y^{(1)}), \dots, (y^{(N)}, x^{(N)}) \right) \quad (1)$$

- where $(x^{(i)}, y^{(i)})$ is one data sample
- e.g (kaggle.com):
 - ① House prices (given location, size, etc.)
 - ② Sales prediction (given advertisement)
 - ③ CO₂ emission (given engine size, etc.)
 - ④ Diabetes (disease progression)
 - ⑤ ...



Linear Regression

- **Goal:** Model the relationship between input and target by a linear equation:

$$\hat{y} = f(x_1, \dots, x_p) = \hat{\beta}_0 + \sum_{j=1}^p \hat{\beta}_j x_j \quad (2)$$

- $\hat{\beta}_0$: Intercept, bias
- $\hat{\beta}_j$: Slope, weight for dimension j
- \hat{y} : Prediction for target y given input x
- Multiple approaches available to fit a linear model e.g.:
 - ① Closed form solution: ordinary least squares
 - ② Locally weighted linear regression
 - ③ Gradient descent
 - ④ ...

Ordinary Least Squares (OLS): Residual Sum of Squares


- **Goal:** find a function f that minimizes residual sum of squares:

$$RSS(\hat{\beta}) = \sum_{i=1}^N \underbrace{\left(y^{(i)} - f(x^{(i)}) \right)}_{\text{residual}}^2 \quad (3)$$

- with

$$f(x_1^{(i)}, \dots, x_p^{(i)}) = \hat{\beta}_0 + \sum_{j=1}^p \hat{\beta}_j x_j^{(i)} \quad (4)$$

- with target value $y^{(i)}$ for input $x^{(i)}$ and model parameters $\hat{\beta} = \{\hat{\beta}_0, \{\hat{\beta}_j\}_{1:p}\}$



Ordinary Least Squares (OLS): Matrix Representations I

- We can replace the sum of $x_j\hat{\beta}_j$ terms (linear combination) with inner product of vector x^T and $\hat{\beta}$:

$$\hat{y} = \hat{\beta}_0 + \sum_{j=1}^p x_j \hat{\beta}_j = \hat{\beta}_0 + \underbrace{x_1 \hat{\beta}_1 + \cdots + x_j \hat{\beta}_j}_{\substack{(x_1 \quad \cdots \quad x_j) \begin{pmatrix} \hat{\beta}_1 \\ \vdots \\ \hat{\beta}_j \end{pmatrix}}} \quad (5)$$

- Which gives us:

$$\hat{y} = \hat{\beta}_0 + x^T \hat{\beta} \quad (6)$$

Ordinary Least Squares (OLS): Matrix Representations II

- Further, we can combine:

$$1\hat{\beta}_0 + \mathbf{x}'^T \hat{\beta} = \begin{pmatrix} 1 & x_1 & \dots & x_j \\ \vdots & & & \end{pmatrix} \begin{pmatrix} \hat{\beta}_0 \\ \hat{\beta}_1 \\ \vdots \\ \hat{\beta}_j \end{pmatrix} \quad (7)$$

- Which gives us the following expression for multiple datapoints:


$$\hat{y} = \mathbf{X}'^T \hat{\beta} \quad (8)$$

- where \mathbf{X}' is a $N \times (p+1)$ matrix containing N feature vectors, $\hat{\beta}$ is a vector containing all parameters and \hat{y} is a vector containing all predictions

Ordinary Least Square (OLS): Closed Form Solution

- Now we can reformulate the RSS objective as:

$$RSS(\hat{\beta}) = (y - \mathbf{X}'^T \hat{\beta})^T (y - \mathbf{X}'^T \hat{\beta}) \quad (9)$$



- By differentiating $RSS(\hat{\beta})$ w.r.t. $\hat{\beta}$, setting the equation to zero, it gives us a unique solution for $\hat{\beta}$ ($\mathbf{X}'^T \mathbf{X}'$ must be invertible):

$$\hat{\beta} = (\mathbf{X}'^T \mathbf{X}')^{-1} \mathbf{X}'^T y \quad (10)$$

Ordinary Least Square (OLS): Solving OLS Manually I

- Consider the following dataset $\mathcal{D} = \{(-1, -2), (1, 0), (2, 1)\}$ where each tuple $(x^{(i)}, y^{(i)})$ consists of a one-dimensional (scalar) feature $x^{(i)}$ and a scalar target $y^{(i)}$. Find the parameters β for a linear regression model using OLS.

$$X = \begin{pmatrix} -1 \\ 1 \\ 2 \end{pmatrix} \quad y = \begin{pmatrix} -2 \\ 0 \\ 1 \end{pmatrix}$$

$$+ \begin{pmatrix} 1 & -1 \\ 1 & 1 \\ 1 & 2 \end{pmatrix} = X'$$

$$AA^{-1} = I$$
$$\begin{pmatrix} 1 & & \\ & 1 & \\ & & 1 \end{pmatrix}$$

$$X' = \begin{pmatrix} 1 & -1 \\ 1 & 1 \\ 1 & 2 \end{pmatrix}$$

$$X^T = \begin{pmatrix} 1 & 1 & 1 \\ -1 & 1 & 2 \end{pmatrix}$$

$$X^T X' = \begin{pmatrix} 3 & 2 \\ 2 & 6 \end{pmatrix} = A = \begin{pmatrix} a & b \\ c & d \end{pmatrix} A^{-1} = \frac{1}{a \cdot d - b \cdot c} \begin{pmatrix} d & -b \\ -c & a \end{pmatrix}$$
$$= \frac{1}{14} \begin{pmatrix} 6 & -2 \\ -2 & 3 \end{pmatrix}$$

Ordinary Least Square (OLS): Solving OLS Manually II


$$X^T y = \begin{pmatrix} -1 \\ 4 \end{pmatrix}$$

$$\begin{aligned} A^{-1} X^T y &= \frac{1}{14} \begin{pmatrix} 6 & -2 \\ -2 & 3 \end{pmatrix} \begin{pmatrix} -1 \\ 4 \end{pmatrix} \\ &= \frac{1}{14} \begin{pmatrix} -14 \\ 14 \end{pmatrix} \end{aligned}$$

$$\hat{\beta} = \begin{pmatrix} -1 \\ 1 \end{pmatrix}$$

Ordinary Least Square (OLS): Solving OLS Manually III

$$\hat{\beta} = \begin{pmatrix} -1 \\ 1 \end{pmatrix}$$



Ordinary Least Square (OLS): Scikit Learn I

`sklearn.linear_model.LinearRegression`



```
class sklearn.linear_model.LinearRegression(*, fit_intercept=True, normalize='deprecated', copy_X=True, n_jobs=None, positive=False)
```

[\[source\]](#)

- `fit_intercept` (bool): Use parameter β_0 during data fit. (Recommended in most use-cases, data must be centered otherwise)
- `normalize`: → will be removed in v1.2
- `copy_X`: if true, X will be copied (otherwise it might be overwritten)
- `n_jobs`: The number of jobs to use for the computation. Requires sufficiently large task.
- `positive`: When set to True, forces the coefficients to be positive.

Going Beyond Linear Functions: Nonlinear Features I

- In reality observed data is often not linear (e.g. quadratic, polynomial, etc.) \rightarrow we still want to use our linear model
- How to extend our previous OLS framework to deal with this kind of data?



Going Beyond Linear Functions: Nonlinear Features II (Polynomial)

- **Single feature x with polynomial target y :**

Suppose: $\mathcal{D} = \{(x^{(j)}, y^{(j)})\}_{1:N}$

$$y(x) = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3$$

- **Polynomial features x with linear target y :**

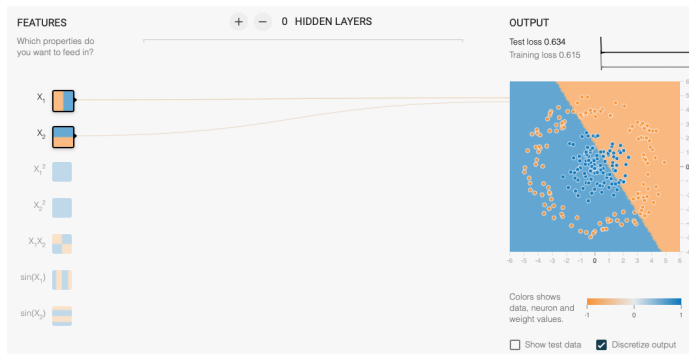
Now: $\mathcal{D} = \{([x^{(j)}, (x^{(j)})^2, (x^{(j)})^3], y^{(j)})\}_{1:N}$

$$y(x) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3$$

- Data is now linear separable in higher dimensional feature space
- General notation: "polynomial feature transform" $\phi(x) = [1, x, x^2, x^3, \dots]$
- Can be expanded to any kind of nonlinear transformation (e.g sin, cosine, sqrt)

Tensorflow Playground¹ - Polynomial Features I

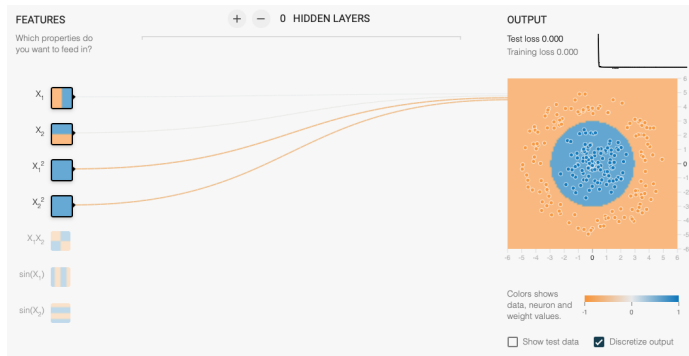
- **Usage:** Taking only input feature without hidden dimension and linear activation function corresponds to linear model optimized using gradient descent.



¹<http://playground.tensorflow.org/>

Tensorflow Playground - Polynomial Features II

- **Summary:** By using nonlinear transformation X_1^2, X_2^2 we can fit our dataset without any modification of the linear model.



Going Beyond Linear Functions: Interactions I

- **Consider:**

$$y = \beta + \beta_1 x_1 + \beta_2 x_2 \quad (11)$$

- **Following:**

$$\frac{\partial y}{\partial x_1} = \beta_1 \quad (12)$$

- This means changing x_1 by one unit will change y by $\beta_1 \rightarrow$ we cannot capture any dependence on x_2 (independent variables)

Exclusive-Or Gate: XOR

Input		Output
x_1	x_2	$x_1 \oplus x_2$
0	0	0
0	1	1
1	0	1
1	1	0

Going Beyond Linear Functions: Interactions II

- **In Contrast:**

$$y = \beta + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_1 X_2 \quad (13)$$

- **Following:**

$$\frac{\partial y}{\partial x_1} = \beta_1 + \beta_3 X_2 \quad (14)$$


- Modeled interactions (combinations) of x_1 with x_2 captures dependencies between those two variables

Exclusive-Or Gate: XOR

Input		Output
x_1	x_2	$x_1 \oplus x_2$
0	0	0
0	1	1
1	0	1
1	1	0

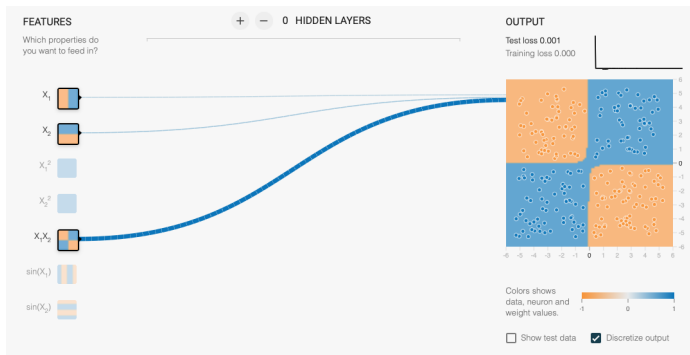
Tensorflow Playground - Interactions I

- **Recap:** Taking only input feature without hidden dimension and linear activation function corresponds to linear model optimized using gradient descent.



Tensorflow Playground - Interactions II

- **Summary:** By using interaction term X_1X_2 we can fit XOR dataset without any modification of the linear model.






`sklearn.preprocessing.PolynomialFeatures`

```
class sklearn.preprocessing.PolynomialFeatures(degree=2, *, interaction_only=False, include_bias=True, order='C') \[source\]
```

- `degree` (int): Maximal degree of the polynomial features (e.g. x_1, x_2 with `degree=2`: $[x_1, x_2, x_1^2, x_2^2, x_1x_2]$)
- `interaction_only` (bool): If True, only interaction features are produced.
- `include_bias` (bool): If True (default), then include a bias column.

Going Beyond Linear Functions: Final Notes



Parametric vs. Nonparametric Learning Algorithms

- **"Parametric" learning algorithm:** fit a fixed set of parameters to data \leftarrow e.g. Linear Regression, Neural Networks, Naive Bayes
- **"Nonparametric" learning algorithm:** amount of data/parameters you need to keep grows (linearly) w.r.t size of data \rightarrow **do not make strong assumptions about the form of the mapping function** (stay flexible). \leftarrow e.g. Locally Weighted Regression,

Going Beyond Linear Functions: Locally Weighted Linear Regression


- Find a set of parameter β' for each point in test set that minimizes weighted sum of squares:

$$WRSS(\hat{\beta}) = \sum_i^m w^{(i)} \left(y^{(i)} - \hat{\beta}^T x^{(i)} \right)^2 \quad (15)$$

- where w is a non-negative weight (scaling factor) with

$$w^{(i)} = \exp \left(-\frac{(x^{(i)} - x)^2}{2\tau^2} \right) \quad (16)$$

- Requires no training phase. Instead coefficients are calculated for each sample during prediction.



Summary: Linear Regression

- Ordinary Least Squares Regression
 - RSS objective and closed form solution for OLS
 - How to solve linear regression problems by hand
- Nonlinear features
 - Polynomial features + how they solve the nonlinearity problem
 - Interactions + how they work
 - Overadaptation of linear models to highly nonlinear features
- Locally Weighted Linear Regression
 - Non-parametric machine learning model
 - Requires no specific assumption about function mapping


Any Questions ?

Recap: Classification

- **Goal:** Find a function f that divides the input space (feature space) X into decision regions separating all classes contained in Y using data \rightarrow discrete outcome:

$$(y^{(1)}, x^{(1)}), \dots, (y^{(N)}, x^{(N)}) \quad (17)$$

- with $y^{(i)} \in \mathbb{N}$ and $x^{(i)} \in \mathbb{R}^n$
- e.g (kaggle.com):
 - 1 Optical character recognition (given image of character)
 - 2 Heart attack classification (given ECG)
 - 3 Skin cancer classification (given skin image)
 - 4 ...



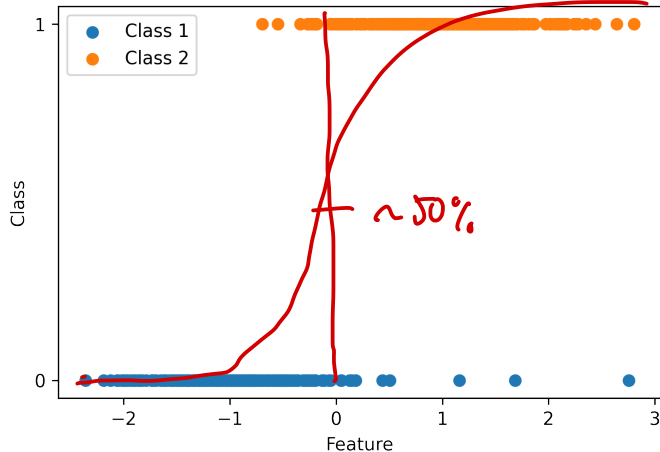
Logistic Regression I: Binary Classification and Class Probabilities

- In contrast to linear regression the type of the dependent (predicted) variable is binary instead of continuous \rightarrow we need some modification to the existing framework.
- Additionally we want to calculate class probabilities using linear functions.
- Given a **binary classification** problem with features $x \in \mathbb{R}^p$, the class probabilities are defined formally as:

$$P(Y = 0|X = x) \tag{18}$$

$$P(Y = 1|X = x) \tag{19}$$


Binary Classification - Features, Target Relationship



Logistic Regression II: Class Probabilities and Odd-Ratio

- Consider ratio between odds of events (odd-ratio):


$$\text{odd-ratio} = \frac{P(Y = 0|X = x)}{P(Y = 1|X = x)} = \frac{P(Y = 0|X = x)}{1 - P(Y = 0|X = x)} \quad (20)$$



- Non-symmetric function
- Odd-ratios for $P \in [0, 1]$ ranges from 0 to infinity

Logistic Regression II: Class Probabilities and Log Odd-Ratio

- Now consider log odd ratio (log-odds)
 $\log(\text{odd-ratio})$
- Symmetric function \rightarrow log-odds against
(-) and in (+) favour while value stays the same



$$P(Y = 0|X = x) \begin{cases} > & P(Y = 1|X = x) : \text{ the log odd-ratio is positive} \\ = & P(Y = 1|X = x) : \text{ the log odd-ratio is zero} \\ < & P(Y = 1|X = x) : \text{ the log odd-ratio is negative} \end{cases} \quad (21)$$

Logistic Regression III: From Log-Odds to Probabilities

- Suppose event with binary outcome with $p = 0.2$ and $\bar{p} = 0.8$ we can get the following conversion rules:

$$odds = \frac{0.2}{0.8} = 0.25 \quad \rightarrow \quad p = \frac{odds}{1 + odds} = 0.2 \quad (22)$$


$$\log(odds) = \ln \frac{0.2}{0.8} = -1.3863 \quad \rightarrow \quad p = \frac{\exp(\log(odds))}{1 + \exp(\log(odds))} = 0.2 \quad (23)$$

- Decision boundaries is defined by the hyperplane where the log-odds are zero
 $\{x | \beta_0 + \beta^T x = 0\}$
- Leaves us with:

$$P(Y = 0 | X = x) = \frac{\exp(\beta_0 + \beta^T x)}{1 + \exp(\beta_0 + \beta^T x)} \quad (24)$$

$$P(Y = 1 | X = x) = \frac{1}{1 + \exp(\beta_0 + \beta^T x)} \quad (25)$$

Logistic Regression VI: Examples for Different Linear Functions



Interactive demo in *logistic_regression_interactive.ipynb*

- **Likelihood:** Factorized probability of a set of N observations as a function of β

$$\mathcal{L}(\beta) = \prod_{i=0}^N P_{\beta}(Y = y_i | X = x_i) \quad (26)$$

- **Maximum Likelihood:** Find the best fitting parameter β that maximizes the likelihood:

$$\hat{\beta} = \arg \max_{\beta} \mathcal{L}(\beta) \quad (27)$$

Logistic Regression VIII: Maximum Likelihood Fit (Log Likelihood)

- Consider the likelihood for $N \rightarrow \infty$ events:

$$\lim_{N \rightarrow \infty} \mathcal{L}(\beta) = \lim_{N \rightarrow \infty} \prod_{i=1}^N P(Y = y_i | X = x_i) \quad \begin{cases} 1 & P(Y = y_i | X = x_i) = 1 \\ 0 & P(Y = y_i | X = x_i) < 1 \end{cases} \quad (28)$$

- Numerical instabilities \rightarrow It is often more convenient to maximize the **log likelihood**
- Log likelihood:**

$$\ell(\beta) = \log \mathcal{L}(\beta) = \sum_{i=1}^N \log P(Y = y_i | X = x_i) \quad (29)$$

Logistic Regression VIII: Maximum Likelihood Fit (Iterative Solutions)


- No closed form solution for logistic regression → instead iterative optimization → convergence to optimal solution not guaranteed.

$$\hat{\beta} = \arg \max_{\beta} \ell(\beta)$$

- or equivalently

$$\hat{\beta} = \arg \min_{\beta} -\ell(\beta)$$

- Many algorithms to find a good maximum likelihood fit: (https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html)



Logistic Regression VIV: Maximum Likelihood - Improved Loss Function

- **Recap:** The loss function of the log-likelihood ℓ is defined as:

$$\ell(\beta) = \log \mathcal{L}(\beta) = \sum_{i=1}^N \log P(Y = y_i | X = x_i) \quad (32)$$

- where y_i is the target class for x_i .

$$P(Y = y_i | X = x_i) = \begin{cases} \log P(Y = 1 | X = x_i) & \text{if } y_i = 1 \\ \log (1 - P(Y = 1 | X = x_i)) & \text{if } y_i = 0 \end{cases} \quad (33)$$

- This gives us :

$$\ell = \sum_{i=1}^N \{y_i \log P(y = 1 | X = x) + (1 - y_i) \log(1 - P(Y = 1 | X = x))\} \quad (34)$$

Logistic Regression X: Dealing with $K \geq 2$ Classes I

- Let $Y = \{ 1, 2, 3, \dots, K \}$
- We want a logistic regression model that provides the posterior probabilities of all K classes.
- **Requirements:**
 - ① Sum to one
 - ② All outputs remain in $[0, 1]$

Logistic Regression X: Dealing with $K \geq 2$ Classes II

- We can define K-1 log-odds using K-1 linear models with:

$$\log \frac{P(Y = 1|X = x)}{P(Y = K|X = x)} = \beta_{10} + \beta_1^T x \quad (35)$$

$$\log \frac{P(Y = 2|X = x)}{P(Y = K|X = x)} = \beta_{20} + \beta_2^T x \quad (36)$$

$$\log \frac{P(Y = K - 1|X = x)}{P(Y = K|X = x)} = \beta_{(K-1)0} + \beta_{K-1}^T x \quad (37)$$

- where the choice of denominator (last class) is arbitrary.

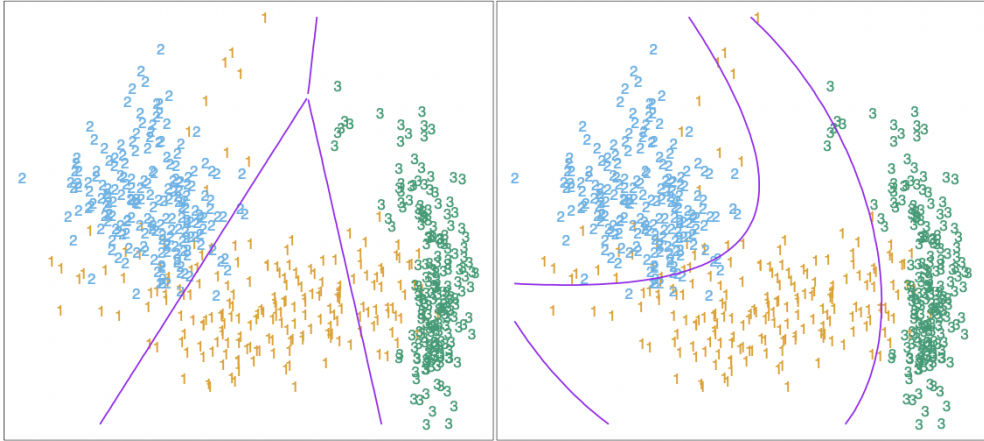
Logistic Regression XI: Dealing with $K \geq 2$ Classes III

- Further we can transform our log-odds to posterior probabilities, respectively as

$$P(Y = 1|X = x) = \frac{\exp(\beta_{k0} + \beta_k^T x)}{1 + \sum_{l=1}^{K-1} \exp(\beta_{l0} + \beta_l^T x)}, k = 1, \dots, K - 1 \quad (38)$$

$$P(Y = K|X = x) = \frac{1}{1 + \sum_{l=1}^{K-1} \exp(\beta_{l0} + \beta_l^T x)} \quad (39)$$

Logistic Regression XI: Dealing with $K \geq 2$ Classes III



Source: Hastie et al.

Notes on Using Linear Regression

- Remember that you can expand your feature space \rightarrow non-linear dependencies
- For nominal features: use one-hot-encoding:
 - Suppose $X \in \{red, green, blue\}$
 - Instead of X use transformed feature: $I(X = x)$
 - red = (1,0,0), green = (0,1,0), blue = (0, 0, 1)

`sklearn.linear_model.LogisticRegression` ¶

```
class sklearn.linear_model.LogisticRegression(penalty='l2', *, dual=False, tol=0.0001, C=1.0, fit_intercept=True,
intercept_scaling=1, class_weight=None, random_state=None, solver='lbfgs', max_iter=100, multi_class='auto', verbose=0,
warm_start=False, n_jobs=None, l1_ratio=None)
```

[\[source\]](#)

- `penalty`: Regularization method (we cover this next lecture).
- `tol`: Tolerance for stopping criteria.
- `fit_intercept`: Specifies if a constant (a.k.a. bias or intercept) should be added to the decision function.
- `class_weight`: Weights associated with classes in the form `class_label: weight`. If not given, all classes are supposed to have weight one.
- `solver`: Algorithm to use in the optimization problem
- `max_iter`: Maximum number of iterations taken for the solvers to converge.

Summary: Logistic Regression

- Relationship between posterior probabilities, odds, and log-odds.
- Modeling of log-odd \rightarrow posterior relationships with linear models.
- Optimization of logistic regression models using maximum likelihood estimation.
- Logistic regression for $K \geq 2$ classes.

Any Questions ?

- ① Hastie, T., Tibshirani, R., Friedman, J. H. (2009). The Elements of Statistical Learning: Data Mining, Inference, and Prediction. Springer.
- ② Deisenroth, M. P., Faisal, A. A., Ong, C. S. (2020). Mathematics for Machine Learning. Cambridge University Press.
- ③ Hosmer, D. W., Lemeshow, S. (2000). Applied logistic regression.
- ④ https://www.montana.edu/rotella/documents/502/Prob_odds_log-odds.pdf

Figures:

- ① https://www.researchgate.net/figure/Schematic-depiction-of-the-locally-weighted-least-squares-kernel-regression-fig3_351477495