

# Modul 215: Machine Learning

## Linear Regression Methods II: Bias, Variance and Regularization

Wintersemester 2023/24

**Tobias Kortus**

**Center for Technology and Transfer (ZTT)**  
University of Applied Sciences Worms

November 16, 2023



# Recap: Linear Regression I


- **Goal:** find a function  $f$  that minimizes residual sum of squares:

$$RSS(\beta) = \sum_{i=1}^N \left( \underbrace{y_i - f(x_i)}_{\text{residual}} \right)^2 \quad (1)$$

- with

$$f(x_1, \dots, x_p) = \beta_0 + \sum_{j=1}^p \beta_j x_j \quad (2)$$

- with  $y_i$  target value for input  $x_i$  and model parameters  $\beta = \{\beta_0, \{\beta_j\}_{1:p}\}$



# Recap: Linear Regression II

- We can reformulate equation (2) as matrix operations:

$$\hat{y} = \mathbf{X}'^T \beta, \quad (3)$$

- where  $\mathbf{X}'$  is a  $N \times (p+1)$  matrix containing  $N$  feature vectors with a leading 1 for the bias term. Now we can reformulate the RSS objective as:

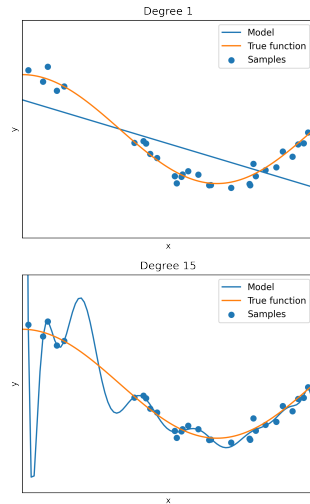
$$RSS(\beta) = \left( y - \mathbf{X}'^T \hat{\beta} \right)^T \left( y - \mathbf{X}'^T \hat{\beta} \right) \quad (4)$$

- By differentiating  $RSS(\beta)$  w.r.t.  $\beta$ , setting the equation to zero, it gives us a unique solution for  $\hat{\beta}$  ( $\mathbf{X}'^T \mathbf{X}'$  must be invertible):

$$\hat{\beta} = \left( \mathbf{X}'^T \mathbf{X}' \right)^{-1} \mathbf{X}'^T y \quad (5)$$

# Recap: Polynomial Feature Transform Interactions

- Transform our features into a higher-dimensional feature space  $\rightarrow$  when selecting good non-linear features, the data might be linear separable.
- **Polynomial Features:**
  - $\phi(x) = [1, x_1, x_2, x_1^2, x_2^2, \dots]$
  - Can be expanded to any kind of nonlinear transformation (e.g. sin, cosine)
- **Interactions:**
  - $\phi(x) = [1, x_1, x_2, x_1x_2, \dots]$
  - Modeled interactions (combinations) of  $x_1$  with  $x_2$  captures dependencies between those two variables



“Numquam ponenda est pluralitas sine necessitate.” (“Plurality must never be posited without necessity”) —William of Ockham

- When you have two competing hypotheses that make the same predictions, the simpler one is most likely the better.
- **Question for this lecture:** How do we quantify the ability of our model? How can we find the best fitting solution with minimum manual effort?

# Recap: Training and Generalization Error

- **Training error:** error (e.g. MSE) calculated on the training set.
- **Test error:** error calculated on the test set.
- **Generalization error:** expectation of our model's error were we to apply it to an infinite stream of additional data examples drawn from the same underlying data distribution as our original sample.
- For MSE we get the following equation for the generalization error:

$$\mathbb{E}_{(x,y) \sim P, D \sim P^n} [(h_D(x) - y)^2] \quad (6)$$

- **Generalization error is a mathematical construct used in statistical learning theory.** We can never calculate the generalization error exactly → we instead estimate the generalization error using the model error on a test set.

# Decomposing the Generalization Error

- We can decompose our generalization error into three main parts:

$$\underbrace{\mathbb{E}_{(x,y,D)} [(h_D(x) - y)^2]}_{\text{Generalization error}} = \underbrace{\mathbb{E}_{x,D} [(h_D(x) - \bar{h}(x))^2]}_{\text{Variance}} + \underbrace{\mathbb{E}_x [(\bar{h}(x) - \bar{y}(x))^2]}_{\text{Bias}^2} \quad (7.1)$$

$$+ \underbrace{\mathbb{E}_x [(\bar{y}(x) - y(x))^2]}_{\text{Noise/ Irreducible Error}} \quad (7.2)$$


- This allows us to further analyze different impacts on the error.

# Decomposing the Generalization Error: Bias Error

- **Recap:** The bias error is defined as:

$$\mathbb{E}_x [(\bar{h}(x) - \bar{y}(x))^2] \quad (8)$$

- **Intuition :** What is the error of the model when trained on infinite data (averaged over all possibilities:  $\bar{h}$ ) given the average over the label (e.g. noisy house prices)






# Decomposing the Generalization Error: Variance Error

- **Recap:** The variance error is defined as:

$$\mathbb{E}_{x,D} [(h_D(x) - \bar{h}(x))^2] \quad (9)$$

- **Intuition :** How much does our classifier change if we train on different data → captures over adaption to data.




# Decomposing the Generalization Error: Noise

- **Recap:** Finally we have noise, defined as:


$$\mathbb{E}_x [(\bar{y}(x) - y(x))^2] \quad (10)$$

- **Intuition :** How big is the data-intrinsic noise?  
This error measures ambiguity due to your data distribution and feature representation. You can never beat this (with your model), it is an aspect of the data.




# Bias Variance Tradeoff

- Increasing the bias will decrease the variance.
- Increasing the variance will decrease the bias.




# Detecting High Bias and High Variance

- **High Variance** (Regime #1)
- Symptoms
  - Training error is much lower than test error
  - Training error is lower than  $\epsilon$
  - Test error is higher than  $\epsilon$
- Remedies
  - Add more training data
  - Reduce model complexity
  - Model ensembles (bagging)





# Detecting High Bias and High Variance



- **High Bias** (Regime #2)
  - Symptoms
    - Training error is higher than  $\epsilon$
    - Test error is higher than  $\epsilon$
- Remedies
  - Use more complex model
  - Add features
  - Model ensembles (boosting)





# Examples of Overfitting: Logistic Regression



# Examples of Overfitting: K-Nearest Neighbor



# Examples of Overfitting: Decision Trees





- One possibility to reduce the variance is to find a subset of the original features that produce the **best results**.
- This also allows for an **easier interpretation** since only a subset of the whole feature set is considered.
- There exist multiple various methods for selecting the best subset of features. We will explore them in the following slides:

# Best-Subset Selection

- Find the best subset of  $k \leq p$  features given an initial set of  $p$  features.
- Use independent validation set/ cross validation for model selection.
- Number of subset sizes  $k$ :  $\binom{p}{k} \approx p^k$
- With efficient algorithms feasible for moderate feature sizes (30-40). [Hastie et al.]
- Not feasible for larger  $p$

# Forward- and Backward-Stepwise Selection

- For  $p \gg 40$  [Hastie et al.] finding the best subset of features becomes feasible  $\rightarrow$  we need good approximations.
- **Forward-Stepwise Selection:**
  - Start with intercept  $\beta_0$
  - Sequentially add new variables to the model that most improve the fit  $\rightarrow$  Quality of algorithm used to identify variables to test defines speed of the approach.
  - Use independent validation set/ cross validation for model selection.
- **Backward-Stepwise Selection:**
  - Start with full model.
  - Sequentially remove variables to the model that has the least impact on the fit.
  - Use independent validation set/ cross validation for model selection.


- Subset selection is a discrete process (variables are either retained or discarded) → often exhibits high variance → doesn't reduce the prediction error of the full model.
- We will explore more continuous alternatives to subset selection.

- **Intuition:** we want to penalize large parameters of the model in order to shrink them towards zero (and each other).
- We therefor penalize the parameters using the sum-of-squares.
- Same concept for neural networks (weight decay, l2-regularization).

$$\hat{\beta}^{ridge} = \arg \min_{\beta} \left\{ \sum_{i=1}^N (y_i - \beta_0 + \sum_{j=1}^p x_j^{(i)} \beta_j)^2 + \lambda \sum_{j=1}^p \beta_j^2 \right\} \quad (11)$$

- where  $\lambda$  is a tunable complexity/ hyperparameter that controls the amount of shrinkage

## Ridge Regression II: A Simple Example



ridge\_lasso\_interactive.ipynb

## Ridge Regression III: Closed Form Solution

- We can rewrite the preceding equation using matrix notation as:

$$RSS(\lambda) = \left(y - \mathbf{X}'^T \hat{\beta}\right)^T \left(y - \mathbf{X}'^T \hat{\beta}\right) + \lambda \beta^T \beta \quad (12)$$

- Similar to the basic OLS we can find a closed form solution:

$$\hat{\beta} = \left(\mathbf{X}'^T \mathbf{X}' + \lambda \mathbf{I}\right)^{-1} \mathbf{X}'^T y \quad (13)$$

- Where  $\mathbf{I}$  is a  $(p + 1) \times (p + 1)$  identity matrix.
- Solution of ridge regression are not equivariant under scaling of the inputs  $\rightarrow$  standardization of features.

# Recap: Standardization

- Standardization is a method used to normalize the range of independent variables or features of data.

$$x' = \frac{x - \bar{x}}{\sigma} \quad (14)$$

- where  $x$  is the original value,  $\bar{x}$  is the mean value of  $x$  and  $\sigma$  is the standard deviation of the value.
- Transformed features have zero mean with unit (1) variance.



# Ridge Regression IV: Limitations

- Ridge regression decreases the complexity of a model but does not reduce the number of variables since it never leads to a coefficient been zero, rather only minimizes it.
- Does not work well for feature reduction.

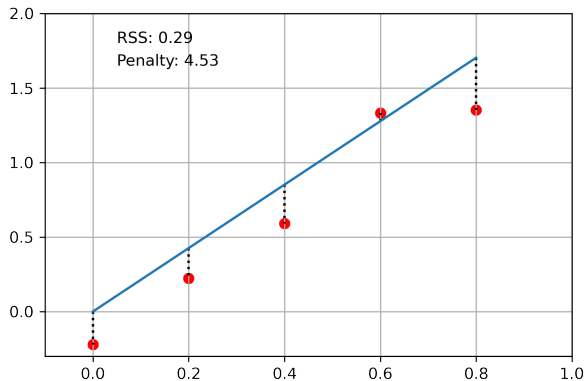
# Lasso Regression

- Similar to ridge regression, we add a penalty term for our parameters  $\beta$ .
- Penalty uses l1 norm  $\sum_j^p |\beta_j|$  ( $|\beta|$ ) instead of l2 norm  $\sum_j^p \beta_j^2$  ( $\|\beta\|$ )

$$\hat{\beta}^{lasso} = \arg \min_{\beta} \left\{ \sum_{i=1}^N (y_i - \beta_0 + \sum_{j=1}^p x_j^{(i)} \beta_j)^2 + \lambda \sum_{j=1}^p |\beta_j| \right\} \quad (15)$$

- where  $\lambda$  is also a tunable complexity/ hyperparameter that controls the amount of shrinkage.
- Same concept is used for neural networks (L1 regularization)
- Lasso constraint makes the solution nonlinear  $\rightarrow$  there is no closed form solution  $\rightarrow$  quadratic programming problem.

# Lasso Regression II: A Simple Example



ridge\_lasso\_interactive.ipynb

# Lasso Regression III: Limitations


- If the number of predictors ( $p$ ) is greater than the number of observations ( $n$ ), Lasso will pick at most  $n$  predictors as non-zero, even if all predictors are relevant.
- If there are two or more highly collinear variables then LASSO regression select one of them randomly which is not good for the interpretation of data

- Combination of Ridge & Lasso regression

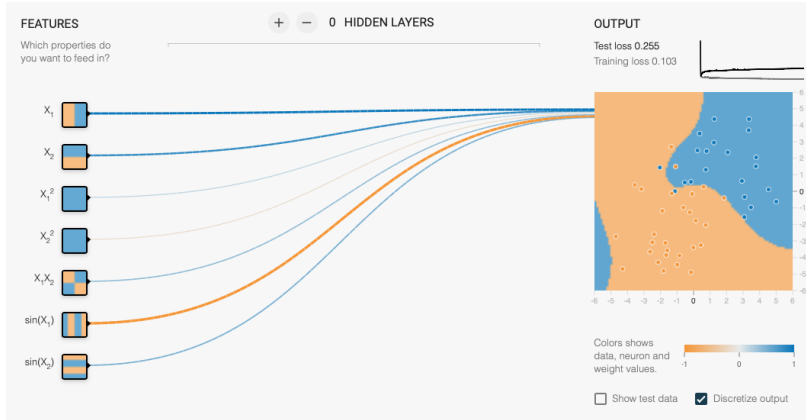
$$\hat{\beta} = \arg \min_{\beta} \left\{ \sum_{i=1}^N (y_i - \beta_0 + \sum_{j=1}^p x_j^{(i)} \beta_j)^2 + \lambda_2 \sum_{j=1}^p \beta_j^2 + \lambda_1 \sum_{j=1}^p |\beta_j| \right\} \quad (16)$$

- where  $\lambda_1$  and  $\lambda_2$  are **individual** regularization factors for ridge and lasso penalty terms.


# Ridge Regression vs. Lasso Regression



# Shrinkage Methods: TensorFlow Playground



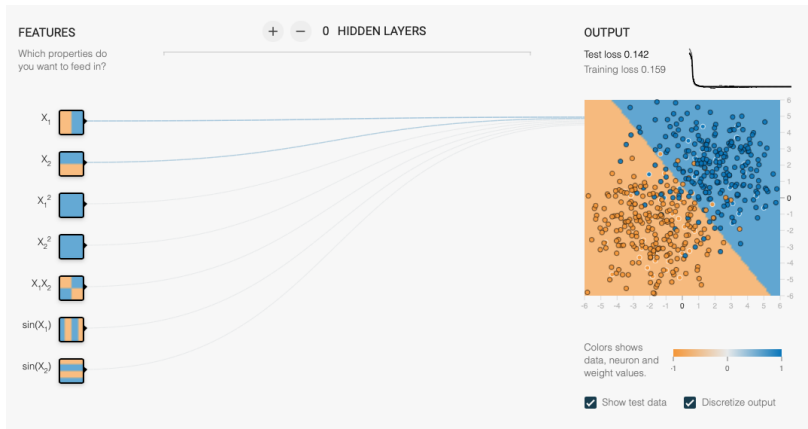
# Shrinkage Methods: TensorFlow Playground



- When enabling test data (checkbox under output) we see that our model overfits to the training data.



# Shrinkage Methods: TensorFlow Playground



- Activating regularization methods (L1/L2) allows us to control the model capacity.

# Controlling Overfitting and Underfitting in Other Models

- Under- and overfitting is also a problem for other machine learning models.
- In most cases, reduce/increase model complexity using model hyperparameters:
- **Decision Tree:**
  - ① max\_depth
  - ② min\_samples\_split
  - ③ min\_samples\_leaf
  - ④ ...
- **K-nearest neighbor:**
  - ① n\_neighbors
  - ② ...
- Read documentation and try modifying parameters (either manual or hyperparameter optimization)

# Final Notes to Linear Models


- We explored multiple linear models. However, there are still many more:
- <https://scikit-learn.org/stable/modules/classes.html>

# (Automatic) Hyperparameter Optimization

- Hyperparameter tuning is just an optimization loop on top of ML model learning to find the set of hyperparameters leading to the lowest error on the validation set.
- Manually optimizing hyperparameters is often not sufficient (requires lots of effort).
- We want to use search algorithms in order to find good/ optimal parameters.
- **Note:** All concepts from model evaluation train, test, eval split; cross-validation; etc. apply also for hyperparameter optimization


# Grid Search

- Exhaustive search over all possible parameter configurations over a user defined grid.
- Requires preliminary knowledge on parameters → specify candidates in search space manually.
- Most straightforward search algorithm that leads to the most accurate predictions as long as sufficient resources are given.



# Random Search

- Randomized search over hyper-parameters from certain distributions over possible parameter values.
- Searching process continues till the predetermined budget is exhausted, or until the desired accuracy is reached
- In most cases, random search is more effective than grid search, but it is still a computationally intensive method.



# Further Optimization Methods

- Bayesian optimization
- Tree parzen estimators
- Evolutionary algorithms
- ...
- Further read:
  - Yu, T., Zhu, H. (2020). Hyper-Parameter Optimization: A Review of Algorithms and Applications. 1-56. <http://arxiv.org/abs/2003.05689>

# A List of Personally Recommended Hyperparameter Optimization Frameworks

- Really simple: [https://scikit-learn.org/stable/modules/grid\\_search.html#](https://scikit-learn.org/stable/modules/grid_search.html#)
- More advanced: <https://optuna.org>



# References

- Hastie, T., Tibshirani, R., Friedman, J. H. (2009). The Elements of Statistical Learning: Data Mining, Inference, and Prediction. Springer.
- Deisenroth, M. P., Faisal, A. A., Ong, C. S. (2020). Mathematics for Machine Learning. Cambridge University Press.
- Yu, T., Zhu, H. (2020). Hyper-Parameter Optimization: A Review of Algorithms and Applications. 1-56. <http://arxiv.org/abs/2003.05689>
- <https://proceedings.neurips.cc/paper/2011/file/86e8f7ab32cfd12577bc2619bc635690-Paper.pdf>
- Bergstra, J., Bardenet, R., Bengio, Y., & Kégl, B. (2011). Algorithms for hyper-parameter optimization. Advances in Neural Information Processing Systems 24: 25th Annual Conference on Neural Information Processing Systems 2011, NIPS 2011, 1-9.
- <https://www.cs.cornell.edu/courses/cs4780/2018fa/lectures/lecturenote12.html>